Quiz for sequence-model

Wednesday, June 26, 2019 11:29 PM

Week 3 Quiz

Quiz, 8 questions

Question 1

1
point

1. Question 1

Why does sequence make a large difference when determining semantics of language?

- [x] Because the order in which words appear dictate their impact on the meaning of the sentence
- [ ] It doesn't
- [ ] Because the order of words doesn't matter
- [ ] Because the order in which words appear dictate their meaning

Question 2

1
point

2. Question 2

How do Recurrent Neural Networks help you understand the impact of sequence on meaning?

- [ ] They don't
- [ ] They look at the whole sentence at a time ❌
- [ ] They shuffle the words evenly
- [x] They carry meaning from one cell to the next

Question 3

1
point

3. Question 3

How does an LSTM help understand meaning when words that qualify each other aren't necessarily beside each other in a sentence?

necessarily beside each other in a sentence?

- [ ] They load all words into a cell state
- [ ] They shuffle the words randomly
- [ ] They don't
- [x] Values from earlier words can be carried to later ones via a cell state

Question 4

1
point

4. Question 4

What keras layer type allows LSTMs to look forward and backward in a sentence?

- [ ] Unilateral
- [ ] Bilateral
- [ ] Bothdirection
- [x] Bidirectional

Question 5

1
point

5. Question 5

What's the output shape of a bidirectional LSTM layer with 64 units?

- [ ] (128,1)
- [x] (None, 128)
- [ ] (128,None)❌
- [ ] (None, 64)❌

Question 6

1
point

6. Question 6

When stacking LSTMs, how do you instruct an LSTM to feed the next one in the sequence?

- [x] Ensure that return_sequences is set to True only on units that feed to another LSTM
- [ ] Do nothing, TensorFlow handles this automatically
- [ ] Ensure that they have the same number of units
- [ ] Ensure that return_sequences is set to True on all units

Question 7

1

1
point

7. Question 7

If a sentence has 120 tokens in it, and a Conv1D with 128 filters with a Kernal size of 5 is passed over it, what's the output shape?

- [x] (None, 116, 128)
- [ ] (None, 120, 128)
- [ ] (None, 120, 124)
- [ ] (None, 116, 124)

Question 8

1
point

8. Question 8

What's the best way to avoid overfitting in NLP datasets?

- [ ] Use LSTMs
- [ ] Use GRUs
- [ ] Use Conv1D
- [x] None of the above